

Fundamentos de Sistemas de Operação

MIEI 2018/2019

Laboratory session 7

Overview

Memory mapped files. Static and dynamic linking.

Testing mmap

Read the manual page of *mmap*. The given *mmcat.c* program copies the contents of a file to the process' standard output, similarly to the *cat* command. Try such program.

```
#include <stdlib.h>
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/mman.h>
#include <sys/stat.h>

void fatal_error(char *str){
    perror(str);
    exit(1);
}

/* mmapcopy - uses mmap to copy file fd to stdout
 */
void mmapcopy(int fd, int size) {
    char *bufp; /* ptr to memory mapped VM area */
    int n;

    bufp = mmap(NULL, size, PROT_READ, MAP_PRIVATE, fd, 0);
    if( bufp == MAP_FAILED) fatal_error("mmap ");

    n = write(1, bufp, size);
    if( n != size) fatal_error("write ");
}

int main(int argc, char *argv[]) {
    struct stat stat;
    int fd;

    /* check for required command line argument */
    if (argc != 2) {
        printf("usage: %s <filename>\n", argv[0]);
        exit(1);
    }

    /* copy the input argument to stdout */
    fd = open(argv[1], O_RDONLY, 0);
    if( fd < 0) fatal_error("error in open");

    if( fstat(fd, &stat) != 0) fatal_error("error in fstat");

    mmapcopy(fd, stat.st_size);
    return 0;
}
```

- Implement a program that copies files (like previous *copy.c*) but without using read/write system calls. You should use *mmap* to transform the copy of files into a copy of memory.
- After testing your program, note that there are files of all sizes, and you can't map all of them to your process memory address space. Then, for the final version of you program, take that into account. Try using as argument to your program the maximum size of your memory buffer, like in previous *copy.c* and *fcopy.c*. Your program should be used like this:

```
mmcopy size file1 file2
```

Notice that:

- the used memory buffer must be a multiple of page size for your architecture. Use `sysconf(_SC_PAGE_SIZE)` to get that page size;
- for changes to the mapped memory of the write be saved, the file to write must be resized to the final size. Use `ftruncate` for that;
- the file must be mapped with `PROT_WRITE` and `MAP_SHARED`;
- and, in order to use the same memory, you must use `munmap` to free each unneeded mapping (similar action can be achieved using option `MAP_FIXED` to enforce the use of the same address). Read `mmap` and `munmap` man pages.

Compare this program performance with *copy.c* and *fcopy.c*, for the same buffer sizes.

Static and dynamic linking

1. Comparing statically- and dynamically- linked executables

Consider any one of the programs used before. Usually, by default, dynamic linking of shared libraries is used by development tools. Make a copy of the current executable file and compile again the program forcing the use of static linking. Example:

```
gcc -static -o mmcat.static mmcat.c
```

Then compare the programs sizes:

```
size mmcat.dynamic
size mmcat.static
```

Try *man ldd* and then execute the following commands to see the used dynamic libraries:

```
ldd mmcat.dynamic
ldd mmcat.static
```

Why the results are different?

Execute each program seeing its use of system calls with `strace`¹ command:

```
strace mmcat.dynamic
strace mmcat.static
```

Verify the differences and observe the *mmap* calls done in the dynamic version by the dynamic linker loader to map *libc* library to the process as needed.

2. Generation of static libraries

Consider the several files that simulates a set of usefull functions that we want to build as a binary code library: *util_file.c*, *util_math.c*, *util_net.c*.

Compile all files as usually and after that build a library by creating an archive file with all the compiled files:

```
cc -c util_file.c util_math.c util_net.c.
ar -rs libmyutil.a util_file.o util_net.o util_math.o
```

To verify the library contents, do:

```
ar -t libmyutil.a
```

and `nm -s libmyutil.a`

¹ If `strace` is not installed use your system package management tool to install it or, if in a debian based distribution, give one of the following commands: `apt install strace`, `apt-get install strace`.

Consider now the program `main.c` that uses your library:

```
#include <stdio.h>
#include "myutil.h"

int main() {
    printf("Inside main()\n");

    /* use a function from each object file that is in the library */
    util_file();
    util_net();
    util_math();
    return 0;
}
```

Compile and link you program using the command:

```
cc -static -o main main.c -L. -lmyutil
```

The option `-L.` gives the current directory as one that contains libraries and `-lmyutil` requests the linking of the `libmyutil` library. If you don't use `"-static"`, your library will still be linked statically but `libc` dynamically. Run the executable confirming that everything works as expected.

To see that `cc` is a *compiler driver* that calls the several phases of the compiler and finally the linker do:

```
cc -v -static -o main main.c -L. -lmyutil
```

3. Generation of dynamic libraries

The generation of a dynamic library is similar to creating a static library. Although, there are two important differences:

1. The compiler must generate code that is position independent. As the processes can load the library in different virtual addresses all the references in the code (jumps, accessing variables, etc) must be relative (to current Instruction Pointer register). In LINUX, this is achieved passing the flag `-fpic` or `-fPIC` to the compiler
2. The tool used to create the library is not `ar`. One must use the `ld` tool (directly or through `cc`). The flag to use is `"-shared"`

You can build both versions of a library but, for now, remove `libmyutil.a`. The sequence of commands used to create a dynamic (also shared) library is:

```
cc -fpic -c util_file.c util_net.c util_math.c
ld -shared -o libmyutil.so util_file.o util_net.o util_math.o
```

or just: `cc -fpic -shared -o libmyutil.so util_file.c util_net.c util_math.c`

Run `file` command to verify the file type of `libmyutil.so`.

Compile and link with your library (you can use `-shared`, but that is the default):

```
cc -o main main.c -L. -lmyutil
```

Try to execute and check the dynamic linking by using `ldd main`. It worked?

At execution time, usually the dynamic loader looks for the shared library file in a pre-defined set of directories (`/lib`, `/usr/lib`, `/usr/X11R6/lib`, ...). One way is to configure the value of the environment variable `LD_LIBRARY_PATH` in order to extend the search to other places. Now do the following:

```
LD_LIBRARY_PATH=. ./main
```

Now it worked? Why? Execute the command

```
LD_LIBRARY_PATH=. ldd main
```

and see why.

To know more about shared and dynamic libraries study the **Program Library HOWTO** from

<http://www.dwheeler.com/program-librar>